AI-Mediated Responsibility:

Structuring, Translating, and Stabilizing Accountability in Sociotechnical Organizations

Manuscript

Sora Kim

Independent Researcher

sorakim.phd@gmail.com

Abstract

Responsibility in regulated organizations does not emerge only after failure; it is continuously produced across anticipatory alignment, procedural embedding, and retrospective stabilization. As artificial intelligence (AI) systems become integrated into documentation and evaluative infrastructures, responsibility is increasingly mediated through computational representation. Yet AI does not create new accountability structures. Instead, it reshapes the temporal dynamics through which responsibility becomes stabilized.

This paper introduces Premature Accountability Convergence (PAC) to describe a conditional acceleration dynamic within the responsibility lifecycle. PAC occurs when computational mediation compresses interpretive space such that responsibility stabilizes more rapidly than causal ambiguity can be meaningfully examined. Under conditions of representational compression, epistemic authority amplification, and temporal contraction of sensemaking, responsibility may converge before alternative structural interpretations remain viable.

By conceptualizing AI as a representational accelerator rather than an autonomous decision-maker, this framework shifts attention from algorithmic bias toward the temporal

Version 1.1 · March 2026

architecture of accountability. PAC extends sociotechnical theories of responsibility by specifying how computational infrastructures intensify stabilization dynamics already embedded within organizational systems. Rather than altering who is responsible, AI may alter how quickly and how narrowly responsibility becomes administratively actionable.

1. Introduction

Artificial intelligence (AI) systems are increasingly embedded within organizational documentation, evaluation, and compliance infrastructures. These systems generate structured representations—summaries, classifications, anomaly flags, predictive indicators—that enter directly into accountability workflows. While frequently described as decision-support tools, AI systems operate at a deeper infrastructural level: they intervene in the representational architecture through which responsibility becomes legible, assessable, and administratively actionable.

Responsibility in regulated organizations is not a singular event but a temporally structured accomplishment. It is anticipatorily aligned before action, procedurally embedded during execution, and retrospectively stabilized under conditions of ambiguity and closure pressure. Prior work within this responsibility program has conceptualized how expectations converge prospectively (Anticipated Accountability Convergence, AAC) and how attribution stabilizes retrospectively through investigative reasoning (Responsibility Convergence, RC). These frameworks explain how responsibility becomes structured and consolidated across stages of organizational life.

What remains theoretically under-specified, however, is the temporal architecture of convergence itself. Existing theories clarify *how* responsibility aligns and stabilizes; they do not fully account for how infrastructural mediation may alter the *tempo* and *breadth* of that stabilization. As AI systems become embedded within documentation and evaluative

processes, interpretive space may narrow earlier in the sensemaking sequence. Convergence may occur more rapidly—not because actors intend closure, nor because institutions become more rigid, but because computational mediation compresses representational diversity and accelerates inscription.

Without a framework that isolates this acceleration dynamic, the effects of AI on accountability risk being misattributed to bias, individual cognition, or institutional design alone. The missing dimension is temporal: how quickly responsibility stabilizes, how much interpretive variance persists before closure, and how infrastructural mediation reshapes these dynamics.

This paper introduces Premature Accountability Convergence (PAC) to conceptualize this acceleration mechanism. PAC identifies the conditions under which computational mediation compresses interpretive space and hastens responsibility stabilization within the responsibility lifecycle. Rather than asking whether AI changes who is responsible, this framework asks how AI reshapes the timing, durability, and scope of responsibility convergence across anticipatory alignment, procedural embedding, and retrospective stabilization. In doing so, it advances a temporal turn in the theory of organizational accountability.

2. Responsibility as a Multi-Stage Organizational Lifecycle

Responsibility in regulated sociotechnical organizations unfolds across temporally distinct yet structurally interconnected stages. Rather than emerging solely at the moment of investigation or formal evaluation, responsibility is progressively produced through anticipatory alignment, procedural embedding, and retrospective stabilization. These stages together constitute a lifecycle through which accountability expectations are generated, operationalized, and ultimately consolidated into administratively actionable forms.

Crucially, responsibility convergence does not occur at a single point. It accumulates across stages. Interpretive variance narrows prospectively, is structured evidentially during execution, and is formally condensed under retrospective closure pressure. Understanding this lifecycle clarifies how accountability is temporally organized prior to the introduction of computational mediation.

2.1 Anticipatory Structuring of Responsibility

In regulated environments, organizational actors rarely act without anticipation of future evaluation. Accountability expectations operate prospectively, shaping how obligations, risks, and acceptable conduct are interpreted before execution. The possibility of inspection, audit, or review functions as a structuring condition of action rather than a subsequent add-on.

Anticipatory alignment occurs when actors orient behavior toward expected evaluative standards. Responsibility is therefore structured ex ante: individuals interpret instructions, exercise discretion, and assess risk with awareness of how these actions may later be scrutinized. Such alignment does not require explicit sanction at the moment of execution; rather, it emerges from embedded expectations that future justification will be required.

At this stage, convergence operates prospectively. Interpretive variance narrows not through formal attribution, but through alignment with anticipated evaluative criteria. Responsibility functions as a precondition for action, organizing conduct around the expectation of later assessability.

2.2 Procedural Embedding During Execution

During execution, responsibility expectations become embedded within documentation systems, reporting architectures, and procedural infrastructures. Standard operating procedures, logging requirements, validation protocols, and compliance frameworks structure how events are recorded and translated into organizational memory.

Procedural embedding stabilizes responsibility not by assigning blame, but by structuring the evidentiary pathways through which future attribution will unfold. Documentation formats define what counts as relevant information, how deviations are categorized, and which explanatory variables become visible within formal records. In doing so, procedural infrastructures pre-format the representational space within which retrospective interpretation will occur.

Convergence at this stage remains incomplete but accumulative. Interpretive possibilities are not eliminated; rather, they are channeled into administratively legible categories. Responsibility is progressively shaped through the structuring of evidence, even before any explicit investigative closure takes place.

2.3 Retrospective Stabilization of Responsibility

After outcomes occur, responsibility is articulated and consolidated through investigations, reports, evaluative documentation, and corrective action processes. Under conditions of causal ambiguity, organizations must translate distributed events into administratively actionable accounts. This requires selecting attributional anchors that can stabilize accountability within formal structures.

Investigative reasoning frequently unfolds under closure pressure: ambiguous causal pathways must be rendered determinate in order to justify corrective action, compliance responses, or performance evaluation. Through this process, interpretive alternatives are gradually reduced as coherent explanatory narratives acquire institutional authority.

Retrospective stabilization marks the formal consolidation of responsibility convergence. Ambiguity is condensed into administratively durable accounts that become embedded in documentation systems and organizational memory. At this stage, responsibility is no longer anticipatory but stabilized; convergence is driven by the need for closure and institutional legibility.

Taken together, these stages illustrate that responsibility convergence is not a discrete event but a cumulative and temporally distributed organizational process.

3. Computational Mediation Within the Responsibility Lifecycle

Having established responsibility as a temporally structured organizational lifecycle, we now consider how computational systems enter and interact with this process. Artificial intelligence systems are increasingly embedded within procedural workflows, documentation regimes, and evaluative infrastructures across regulated domains. Rather than functioning as autonomous decision-makers, these systems operate as representational mediators that translate organizational events into structured, recordable, and evaluable forms.

Computational outputs—such as summaries, classifications, anomaly flags, pattern detections, and predictive indicators—do not independently assign responsibility. However, they shape how events are rendered legible within organizational documentation and evaluative processes. By structuring information into standardized and repeatable representations, computational mediation influences how actors interpret, record, and retrospectively reconstruct events.

Importantly, computational mediation does not introduce responsibility structures ex nihilo. The anticipatory alignment, procedural embedding, and retrospective stabilization of responsibility precede the integration of AI systems. Organizational infrastructures already generate expectations of accountability and stabilize attribution through investigative processes. What computational mediation alters is not the existence of responsibility convergence but its tempo and representational form.

Because algorithmically generated representations often carry the perceived authority of formalized calculation, they may narrow interpretive divergence earlier in the sensemaking process. Structured outputs can foreground particular causal pathways while backgrounding

others, shaping which interpretations appear administratively actionable. In doing so, computational systems may compress the temporal space between ambiguity and stabilization within the responsibility lifecycle.

This interaction introduces a distinct dynamic: responsibility expectations may align and stabilize more rapidly when mediated through computationally structured representations. The lifecycle remains intact, but its stabilization phase may be accelerated through representational compression.

4. Premature Accountability Convergence (PAC)

Computational mediation does not introduce responsibility structures into organizational systems; rather, it interacts with a pre-existing lifecycle of convergence. Within this lifecycle—anticipatory structuring, procedural embedding, and retrospective stabilization—AI systems function as representational accelerators. They do not determine who is responsible, but they may alter the tempo and breadth through which responsibility stabilizes.

Premature Accountability Convergence (PAC) identifies a structural acceleration tendency that becomes visible under conditions of computational mediation. PAC occurs when representational compression, perceived epistemic authority, and temporal contraction combine to reduce interpretive variance before causal ambiguity has been meaningfully examined. The lifecycle itself remains intact; what shifts is the timing and scope of stabilization.

PAC is not a claim that AI inherently distorts accountability. Rather, it specifies a conditional infrastructural dynamic: when structured computational outputs enter documentation and evaluative systems early in the sensemaking sequence, responsibility may converge more rapidly and with diminished interpretive plurality. Convergence, in such cases, is not imposed but accelerated.

PAC therefore reframes computational mediation as a temporal intervention within accountability architecture. It draws attention to how infrastructural design shapes the durability of ambiguity, the persistence of alternative explanations, and the pace at which attribution becomes administratively actionable.

Three interrelated mechanisms contribute to PAC: representational compression, epistemic authority amplification, and temporal contraction of sensemaking.

4.1 Representational Compression

Computational systems translate complex and distributed organizational events into structured outputs—summaries, classifications, risk scores, anomaly flags, and predictive indicators. In doing so, they compress interpretive multiplicity into standardized representational forms.

This compression enhances administrative clarity and traceability. However, it also pre-selects which explanatory variables remain visible within formal documentation. As events are formatted into predefined categories, the range of plausible causal accounts may narrow prior to full deliberation. Interpretive variance does not disappear, but it becomes institutionally backgrounded.

Representational compression thus operates as the primary structural driver of PAC.

4.2 Epistemic Authority Amplification

Algorithmically generated outputs often carry perceived legitimacy derived from formalized calculation, technical infrastructure, and data-driven framing. When computational representations are treated as epistemically authoritative, interpretive contestation may diminish.

Epistemic authority amplification does not require blind trust in algorithms. Rather, it emerges when structured outputs become embedded within formal documentation systems and are granted institutional weight. Under conditions of ambiguity, such outputs may anchor

explanatory frames earlier in the stabilization process, accelerating convergence around computationally foregrounded interpretations.

4.3 Temporal Contraction of Sensemaking

Computational mediation frequently reduces the interval between event occurrence and formal inscription. Automated logging systems, real-time dashboards, and immediate classification tools shorten the deliberative window within which alternative interpretations may circulate informally.

When documentation is generated rapidly and embedded early within evaluative infrastructures, preliminary interpretations acquire durability before distributed explanations are fully articulated. Responsibility may therefore stabilize prior to the exhaustion of interpretive diversity.

Temporal contraction amplifies the acceleration dynamic initiated by representational compression.

Representational compression constitutes the primary structural driver of PAC, translating distributed organizational events into administratively legible forms. Yet acceleration does not operate uniformly. Its intensity depends on contextual amplification mechanisms—most notably the perceived epistemic authority granted to computational outputs and the temporal interval available for interpretive deliberation. When these conditions converge, responsibility stabilization may occur earlier and with diminished interpretive plurality. The following propositions formalize these conditional dynamics.

5. Theoretical Propositions: Premature Accountability Convergence

Proposition 1. Representational Compression

The greater the degree of computational representational compression within organizational documentation systems, the more likely interpretive space will narrow prior to full causal

exploration.

Computational representational compression refers to the extent to which complex, distributed, and ambiguous events are translated into standardized, categorical, or summarized formats. As organizational documentation systems rely more heavily on structured computational outputs—such as classification schemas, risk categories, or summary dashboards—interpretive possibilities become formatted into predefined representational containers. While such structuring enhances administrative clarity and efficiency, it simultaneously limits the range of causal alternatives that remain visible within formal documentation. When ambiguity is compressed into categorical outputs, alternative explanatory pathways may lose salience before being fully explored, thereby increasing the likelihood of premature convergence.

Proposition 2. Epistemic Authority

The stronger the perceived epistemic authority of computational outputs, the more rapidly responsibility expectations will align around computationally foregrounded interpretations.

Epistemic authority amplification occurs when algorithmically generated outputs are treated as objective, neutral, or technically superior forms of knowledge. Because computational representations derive legitimacy from formalized calculation and technical infrastructure, actors may defer to their outputs with reduced contestation. In contexts of ambiguity, perceived data-driven precision can accelerate interpretive alignment by stabilizing particular explanatory frames early in the sensemaking process. As actors converge around computationally foregrounded interpretations, alternative causal accounts may receive diminished scrutiny, thereby contributing to accelerated stabilization within the responsibility lifecycle.

Proposition 3. Temporal Contraction

The shorter the temporal interval between event occurrence and formal computational

documentation, the greater the likelihood of premature responsibility stabilization.

Temporal contraction of sensemaking refers to the reduction in deliberative interval between an event and its formal inscription within documentation systems. Real-time dashboards, automated incident logging, and immediate classification tools compress the period during which interpretive divergence may be explored informally. When documentation is generated rapidly and becomes embedded within evaluative infrastructures, early interpretations gain durability and institutional visibility. As a result, responsibility accounts may solidify before distributed or systemic explanations are meaningfully articulated, increasing the probability of premature convergence.

Taken together, these propositions specify PAC as a conditional intensification dynamic within the responsibility lifecycle. Computational mediation does not create responsibility convergence ex nihilo; rather, it may accelerate stabilization and compress interpretive diversity under specific infrastructural conditions. Recognizing PAC enables more precise theoretical analysis of how accountability becomes structured and consolidated in AI-supported organizational environments.

6. Discussion: AI Mediation within the Responsibility Lifecycle

PAC contributes to sociotechnical and organizational theories of accountability by introducing acceleration as a critical analytical dimension within responsibility structuring. Prior research has examined how responsibility is anticipatorily aligned before action (AAC) and how attribution stabilizes after events through investigative processes (RC). However, less attention has been paid to how computational infrastructures alter the tempo and breadth of stabilization across these stages. By conceptualizing AI as a representational accelerator rather than an autonomous decision-maker, PAC extends accountability theory beyond attribution and compliance toward the temporal dynamics of responsibility production.

Within the broader responsibility lifecycle, responsibility is not a singular event but a staged organizational accomplishment. It is anticipatorily structured through procedural and evaluative expectations, embedded during execution through documentation architectures, stabilized retrospectively through investigative closure, and increasingly mediated through computational representation. PAC does not replace these stages; rather, it operates across them. By compressing interpretive possibilities, amplifying epistemic authority, and contracting deliberative time, computational mediation may intensify convergence dynamics that already exist within accountability infrastructures.

The central tension exposed by PAC is therefore not between automation and human judgment, but between speed and deliberation. As computational systems enhance administrative efficiency, they may simultaneously reduce the durability of ambiguity within organizational processes. Accountability systems differ not only in who they assign responsibility to, but in how long interpretive variance is permitted to endure before stabilization becomes institutionally binding. When interpretive space contracts too rapidly, distributed or structural explanations may lose durability before being meaningfully examined.

This perspective reframes AI not as an external disruptor of responsibility, but as an infrastructural layer that interacts with existing stabilization mechanisms. AI systems do not inherently produce new responsibility structures; instead, they reshape the timing and scope through which responsibility becomes administratively actionable. The central theoretical implication is therefore not technological determinism, but infrastructural acceleration. Responsibility convergence remains contingent on organizational ambiguity, closure pressure, and accountability design. Computational mediation may heighten these dynamics, particularly when representational outputs are granted authority and rapidly embedded in formal documentation.

Situating PAC within the responsibility program clarifies its relationship to adjacent frameworks. Anticipated accountability convergence (AAC) explains how expectations align before action through procedural and documentary infrastructures. Responsibility convergence (RC) explains how attribution stabilizes after events under investigative closure pressure. PAC identifies how computational mediation may accelerate stabilization dynamics across these stages. Rather than constituting a separate theory of accountability, PAC articulates a cross-stage accelerative mechanism embedded within contemporary sociotechnical systems.

In this sense, accountability should be understood not merely as a normative principle, but as an architecture of temporal structuring and evidentiary design. This integration suggests that accountability infrastructures must be analyzed not only in terms of fairness, transparency, or bias, but also in terms of temporal architecture. When computational systems compress interpretive space too rapidly, organizations may inadvertently reduce the durability of distributed explanations.

Preserving interpretive plurality may therefore require deliberate design interventions, including mechanisms for contestability, temporal buffering, layered documentation that preserves alternative interpretations, and procedural conditions under which stabilized accounts can be reopened. Recognizing PAC as an infrastructural acceleration dynamic enables more precise analysis of how accountability architectures evolve under computational mediation.

PAC thus foregrounds the temporal politics of responsibility stabilization. In AI-supported environments, responsibility may converge not because computational systems impose determinate judgments, but because they reshape the timing, authority, and durability of representational translation. Attending to this acceleration dynamic opens a pathway toward designing accountability systems that balance administrative efficiency with sustained

deliberative capacity.

7. Conclusion

This paper introduced Premature Accountability Convergence (PAC) as a conditional acceleration dynamic within contemporary accountability infrastructures. Rather than portraying AI as a determinant of responsibility, the framework conceptualized computational systems as representational mediators that may compress interpretive space and intensify stabilization dynamics within the responsibility lifecycle. PAC therefore reframes AI's role in organizations not as a source of new responsibility structures, but as an infrastructural force that alters the timing and scope of convergence.

By situating PAC within a broader responsibility program—including anticipatory convergence, procedural embedding, and investigative stabilization—this analysis demonstrates that responsibility in regulated organizations is temporally produced across multiple stages. Computational mediation interacts with these stages by accelerating representational consolidation, particularly under conditions of ambiguity and epistemic authority amplification.

Recognizing this acceleration dynamic enables more precise theoretical and empirical investigation of accountability design in AI-supported environments. Rather than asking whether AI improves or undermines responsibility, future research may examine how organizational systems structure interpretive time, contestability, and representational plurality in order to preserve meaningful deliberation within accountability processes. PAC thus offers a conceptual foundation for analyzing how computational infrastructures reshape the temporal architecture of responsibility in contemporary sociotechnical organizations.

References

Bovens, M. (2007). Analysing and assessing accountability: A conceptual framework. *European Law Journal, 13*(4), 447–468. https://doi.org/10.1111/j.1468-0386.2007.00378.x

Porter, T. M. (1995). *Trust in numbers: The pursuit of objectivity in science and public life*. Princeton University Press.

Power, M. (1997). *The audit society: Rituals of verification*. Oxford University Press.

Suchman, M. C. (1995). Managing legitimacy: Strategic and institutional approaches. *Academy of Management Review, 20*(3), 571–610. https://doi.org/10.5465/amr.1995.9508080331

Weick, K. E. (1995). *Sensemaking in organizations*. Sage Publications.